🧾 Parchi

Hinglish voice notes → verified orders → spoken replies · Gnani AI Great Indian AI Internship Challenge 2026
Built on **Gnani Prisma v3.0** (speech-to-text) and **Gnani Timbre v2.5** (text-to-speech)

🎥 Demo video (1:38): github.com/arnav-156/Parchi/blob/main/docs/Parchi_Demo.mp4 · 💻 Code: github.com/arnav-156/Parchi

PROBLEM

India's independent FMCG distributors each serve 100–300 kirana stores. Retailers order by WhatsApp voice note, in Hinglish: "*Parle-G ke do carton, Maggi chhota wala paanch dabba…*". Staff listen to every note and retype each line into Tally. Retailers rarely get a confirmation, and every slip becomes a wrong delivery.

SOLUTION

Parchi is an order desk that turns those notes into a verified order sheet, and **talks back to the retailer** in their language.

1. **Hear**: the note is transcribed by **Gnani Prisma v3.0** (`indic-auto`, `codemix`), biased with the distributor's own catalog as keywords, so brand names come out right.
2. **Understand**: deterministic rules parse Hindi/English numbers, fractions (*dedh, saadhe*) and trade units (*peti, dabba, darjan*), then match each line to a SKU. Complaints ("*pichli baar kam aaya*") become notes.
3. **Flag, don't guess**: GREEN confident · AMBER needs one answer (which pack? how many?) · RED not in catalog.
4. **Ask back**: for amber lines, **Gnani Timbre v2.5** speaks a Hinglish question ("*Colgate kaunsa chahiye, sau gram ya do sau gram?*"). The retailer's voice reply resolves the line.
5. **Confirm & export**: a human approves (never automatic), the retailer gets a spoken confirmation, and the distributor gets a Tally-ready CSV.

RESULTS (synthetic test set, 59 order lines)

	Without catalog biasing	With Parchi biasing
Line F1, phone-quality audio	0.800	**0.820**
Product (SKU) accuracy, phone audio	94.9%	**100%**
Wrong lines auto-marked green	4.3%	**3.8%**

End-to-end browser test with synthetic phone calls: **16/16 pass, 0 wrong lines approved**. Prisma returns a transcript in ~0.6 s per note. 532 automated tests; every Gnani call is cached and counted.

WHAT'S REAL VS MOCKED

Real: Prisma v3.0 transcription, Timbre v2.5 speech, parsing, matching, review, clarification loop, export.
Mocked: WhatsApp (file upload) and Tally (CSV). All data is synthetic, with no real customers or recordings.

Next: WhatsApp Business API, direct Tally push, and a held-out test set of real retailer recordings.